

Return Path - R Model

October 21, 2016

The R Model builds on the test and train files from the Python data pipeline. A Random Forest model is trained with 500 trees and a depth of 3.

Ideally, a cross validation should be done, however the initial model seems to work with good accuracy.

Following is the code.

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.2.5
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
## Build a Random Forrest Model
setwd("C:/Users/gnagaraj/Downloads")

#Prepare the Test and Train Datasets
train_df=read.csv("Train_split.csv",colClasses = c("Discretized_Read_Rate"="factor"))
test_df=read.csv("Test_split.csv",colClasses = c("Discretized_Read_Rate"="factor"))

#Drop unwanted columns
train_df <- subset(train_df,select = -c(X,id,read_rate))
test_df <- subset(test_df,select = -c(X,id,read_rate))

#Train a Random Forrest Model with 500 trees and depth of 3
rf <- randomForest(Discretized_Read_Rate ~ .,data = train_df)

#Predict the test data
predictions<-predict(rf,subset(test_df,select = -c(Discretized_Read_Rate)))

#Calculate the Confusion Matrix
cm<-table(as.numeric(unlist(test_df['Discretized_Read_Rate'])),as.numeric(predictions
))
print(cm)
```

```
##
##      1      2      4      5      6      7      8      9     10     11     12     13     14
##  1  2627      0      0      0      0     86     17     10      5      1      1      0      0
##  2      6      0      0      6      0      1      0      4      0      3      2      1      0
##  3      6      0      1      7      0      1      0      1      1      0      2      0      0
##  4      1      0      0     11      1      0      1      0      2      0      2      0      0
##  5      2      0      1    189      1      0      0      0      2      1      3      0      0
##  6      4      0      0     32     19      0      0      1      1      0      0      0      0
##  7    246      0      1      0      0    174     18     19      2      2      0      0      0
##  8    122      0      0      0      0     83     80     15      6      4      5      0      0
##  9     48      0      0      2      0     36     35     35      7      2      1      0      0
## 10     34      1      0      3      0     11     27     13     22      8      3      0      0
## 11     25      0      0      4      0      8      7     14      5     11      5      0      0
## 12     12      0      0      3      0      3      6      5     11     10     12      1      0
## 13     12      0      2      7      0      3      1      6      2      9      4      1      0
## 14      4      0      0      8      0      1      2      4      2      1      2      1      1
```

```
#Print the Accuracy
print(c("Accuracy",sum(diag(cm)/sum(cm))))
```

```
## [1] "Accuracy"          "0.620308542482155"
```